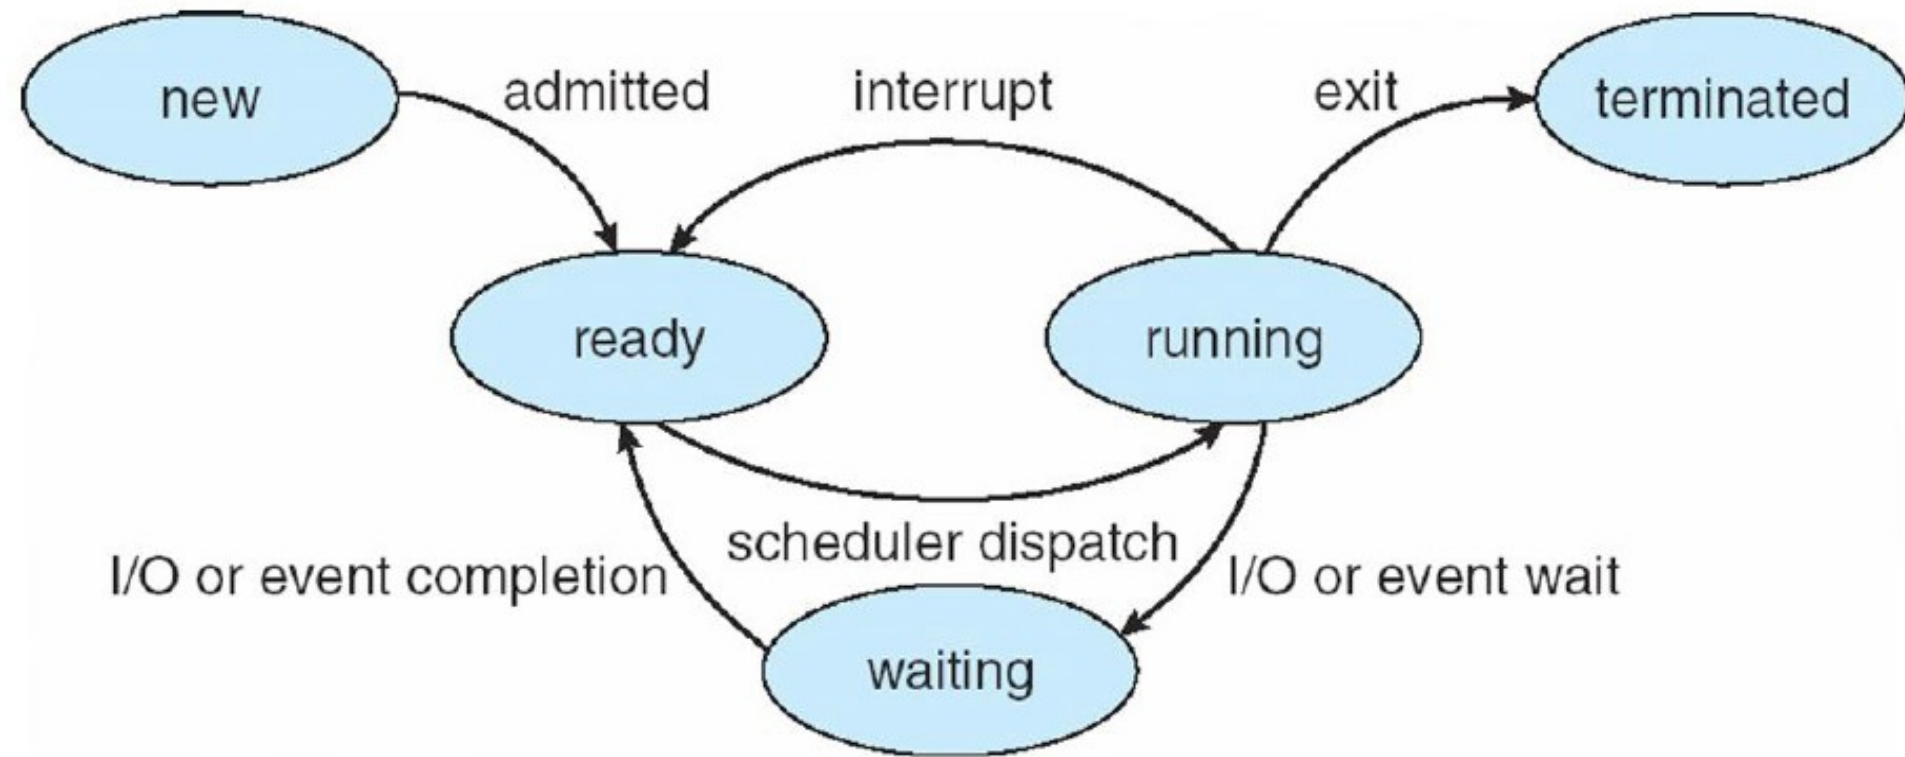


Linux Scheduler

OS-323 (Spring 2013)

Process states



CPU scheduler

- Makes the computer more productive by switching the CPU among processes
- CPU scheduling may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from new or waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- The other scheduling is **preemptive**

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority
- Problem: Starvation – low priority processes may never execute
- Solution: Ageing – as time progresses increase the priority of the ready processes

Round Robin (RR)

- Each process gets allocated a small unit of CPU time (time quantum)
- After this time has elapsed, the process is preempted
- With N processes in the ready queue and the time quantum of q , no process waits more than $(N-1)q$ time units.
- q must be large with respect to context switch, otherwise overhead is too high

Scheduling in Linux

- Initially a circular queue with a round-robin scheduling policy
 - Efficient adding and removing of processes
 - Simple and fast for a small number of tasks
- Linux 2.4: $O(N)$ scheduler
- Linux 2.6: $O(1)$ scheduler
- Linux 2.6: Completely Fair Scheduler (CFS)

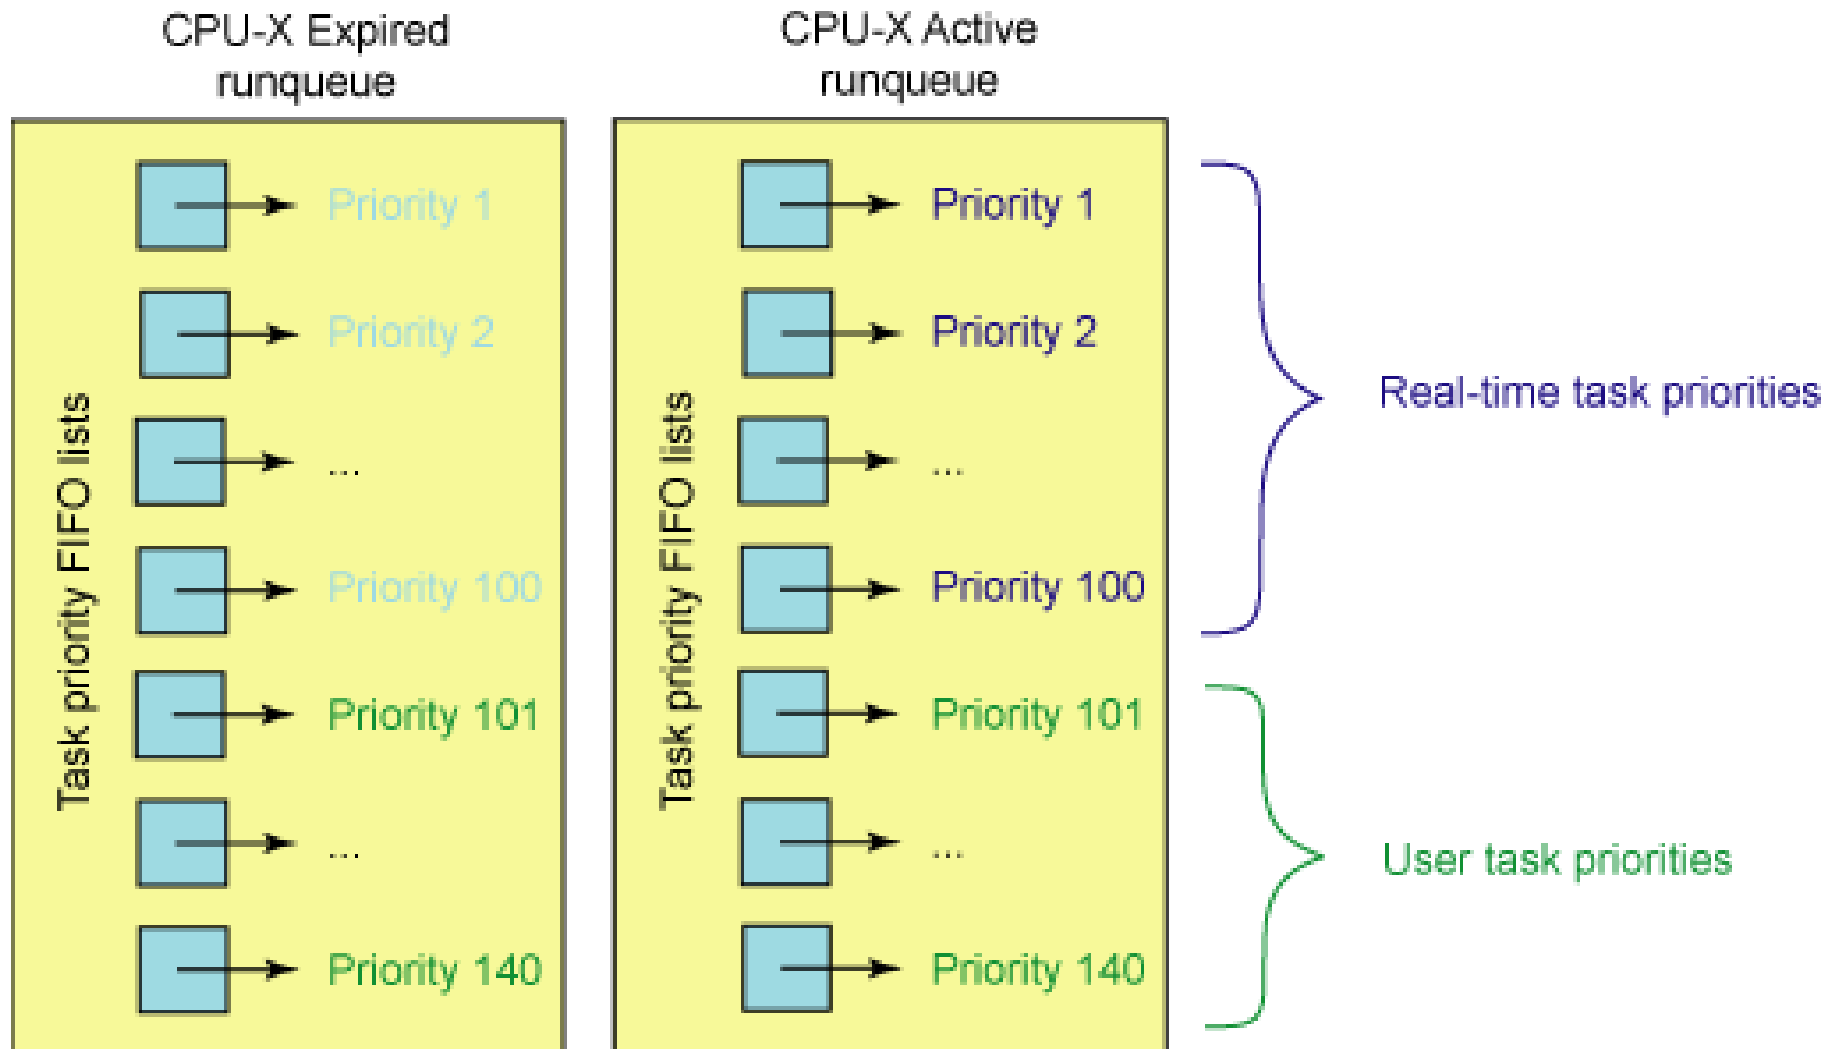
O(N) Scheduler

- Multilevel feedback queue with 140 priority levels
- CPU time divided into epochs
- Within each epoch, every task was allowed to execute up to its time slice
- During each scheduling event, the scheduler would iterate over all tasks, applying a *goodness* function (metric) to determine which task to execute next

O(1) Scheduler

- Scheduling done completely in $O(1)$
- Runqueue consists of two arrays – active and expired
- Each array consists of 140 linked lists, that are serviced in FIFO order, each representing a priority level
- The first 100 levels are reserved for real-time tasks, and the remaining 40 for user tasks
- Each task is assigned a timeslice based on its priority

O(1) Scheduler



O(1) Scheduler

- Enqueue: add to the end of one of the fixed number of lists
- Dequeue: deletion from a linked list is const. time
- Pick next task: choose the first task in the highest priority list that is not empty. A bitmap with one bit per list is used to quickly determine the first nonempty list.
- Preventing starvation: When a task on the active array uses up its timeslice, it is moved to the expired array. Once all tasks move to the expired array, arrays are swapped.
- All scheduling operations depend only on the number of priority levels, and not the number of tasks

O(1) Scheduler

- SMP support:
 - Per-CPU runqueues: better cache utilization
 - Big-lock architecture replaced with a lock per runqueue
 - Task migration for load balancing

O(1) Scheduler

- Task priorities are dynamically altered during execution (and accordingly timeslices)
- Goal: prioritize interactive tasks over batch tasks
- Interactive tasks typically spend most time sleeping, they have short bursts of CPU activity, before blocking waiting for user input
- A heuristic is used to determine whether a task is interactive based on the ratio of sleep time to execution time

CFS

- Motivated by Rotating Staircase Deadline Scheduler by Con Kolivas which showed that a simple fair scheduler performed better than $O(1)$ scheduler's interactivity heuristics
- Main idea: Allocate the fair share of CPU time to each of the ready tasks. If there are N ready tasks, each should get $1/N$ of CPU power.

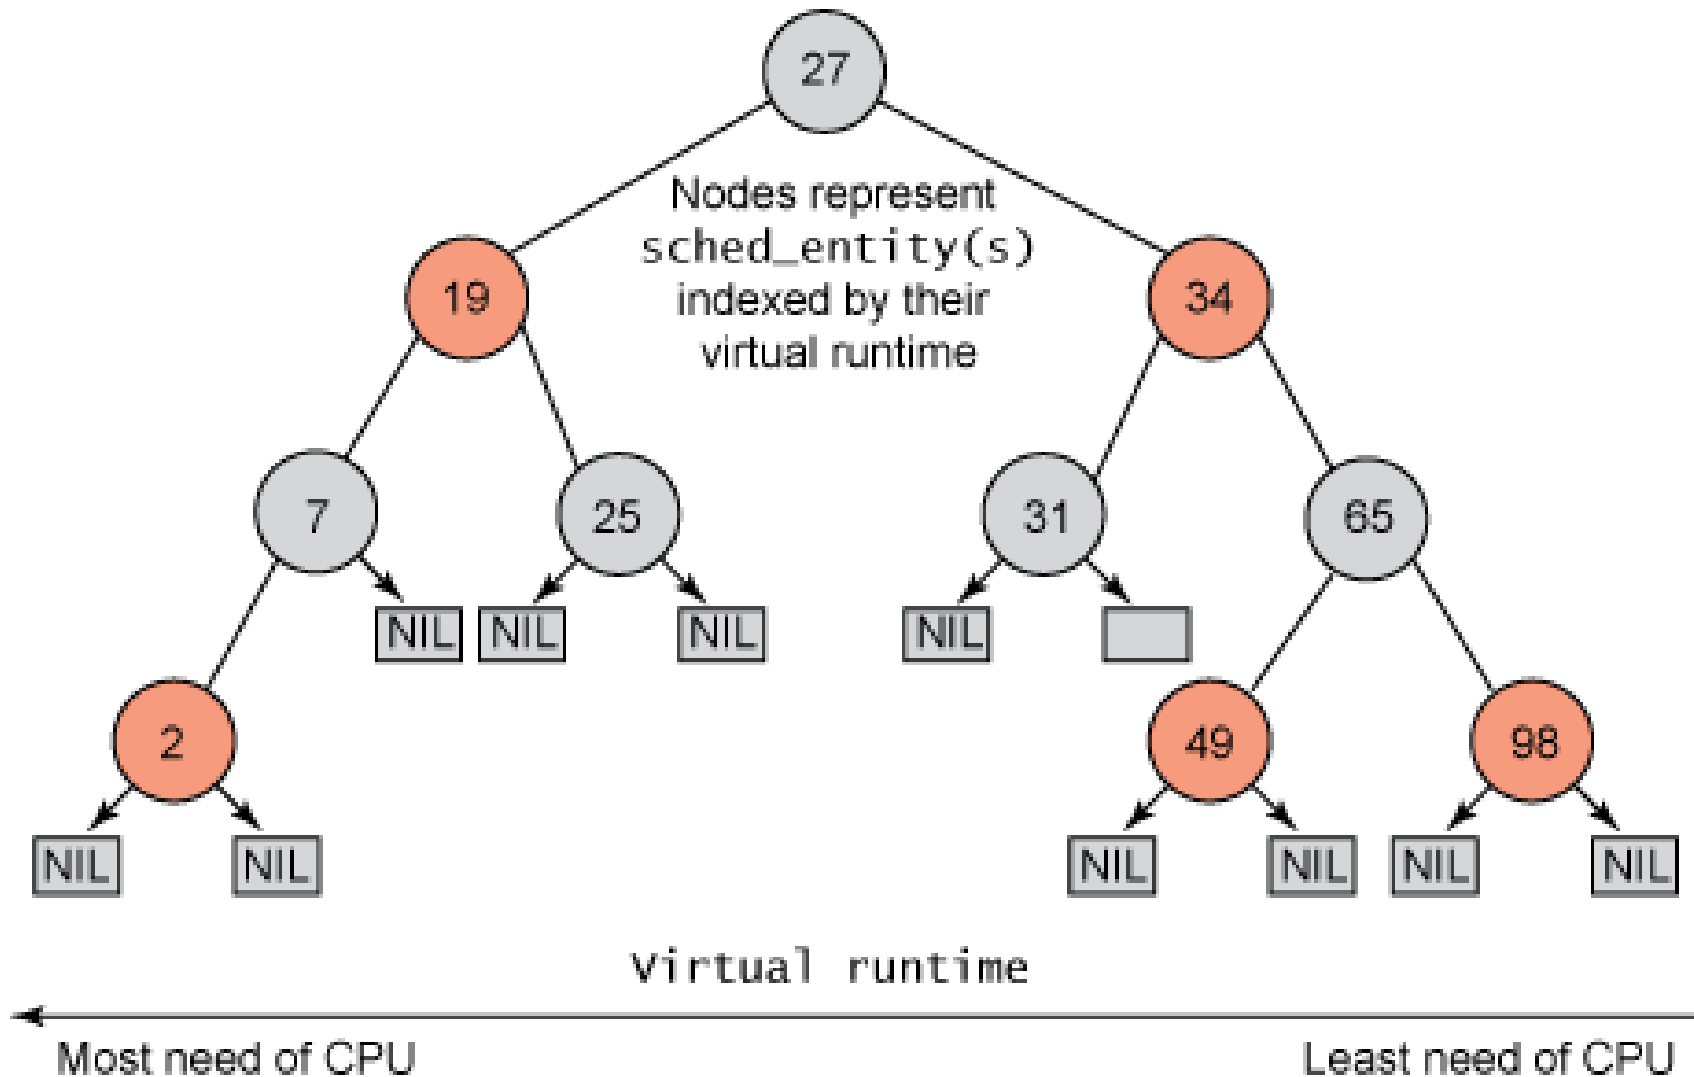
CFS

- Each task has a virtual runtime value, which is its actual runtime normalized to the number of ready tasks
- Task priority is incorporated as a weight factor into this formula
- The CPU is allocated to the task with the smallest virtual runtime value
- The concept of timeslices is removed from CFS
- Instead, there is a scheduling granularity parameter (in ns), that defines how often a new task is chosen

CFS

- The runqueue in CFS is replaced by a red-black tree
- A red-black tree is a self-balancing binary search tree. It has guaranteed $O(\log N)$ complexity of insertion, deletion and rebalancing.
- The red-black tree is ordered by task virtual runtime
- The leftmost leaf of the tree always has the smallest virtual runtime and is chosen for scheduling in $O(1)$
- During execution, its virtual runtime increases and it moves to the right in the tree, so that the next task with the smallest virtual runtime comes to the leftmost position
- Whenever the running task's CPU usage is accounted for, it is reinserted in the tree in $O(\log N)$ time

CFS



CFS

```
struct cfs_rq{
    struct load_weight load;
    unsigned int nr_running;

    u64 min_vruntime;

    struct rb_root tasks_timeline;
    struct rb_node *rb_leftmost;

    struct sched_entity *curr;
}
```

RT Scheduler

- Schedules tasks for 100 highest priority levels
- The first task in the highest nonempty priority list is chosen
- FIFO policy: The running task executes without preemption
- Round robin policy: The tasks have timeslices, and are preempted and put at the end of the list upon timeslice expiry

Priorities

- Priorities
 - [0-99] are kernel RT priorities
 - [100-139] are kernel user priorities
 - User priorities [100-139] correspond to nice priorities [-20-19]
 - Default priority is kernel 120 (or nice 0)
- `static_prio`: static priority, can be changed only through system calls (`setpriority`, `nice`)
- `normal_prio`: inverse of static priority used in RT scheduler
- `prio`: dynamic priority

Scheduling classes

- The CFS scheduler also introduced scheduling classes, an extensible hierarchy of scheduling modules
- Each class is implemented with a `sched_class` structure that contains hooks to functions to be called when some interesting events occur

Scheduling classes

```
struct sched_class {
    const struct sched_class *next;

    void (*enqueue_task) (struct rq *rq, struct task_struct *p, int flags);
    void (*dequeue_task) (struct rq *rq, struct task_struct *p, int flags);

    void (*yield_task) (struct rq *rq);
    void (*check_preempt_curr) (struct rq *rq, struct task_struct *p, int flags);

    struct task_struct * (*pick_next_task) (struct rq *rq);
    void (*put_prev_task) (struct rq *rq, struct task_struct *p);
    void (*set_curr_task) (struct rq *rq);

    void (*task_tick) (struct rq *rq, struct task_struct *p, int queued);

    void (*switched_from) (struct rq *this_rq, struct task_struct *task);
    void (*switched_to) (struct rq *this_rq, struct task_struct *task);
    void (*prio_changed) (struct rq *this_rq, struct task_struct *task, int oldprio);
};
```

Scheduling classes

- enqueue_task/dequeue_task
 - Called when a task enters/leaves a ready state
 - Activate/deactivate (new task/terminataion, wakeup/sleep)
 - Changed priority (rt_mutex_setprio, set_user_nice)
 - Changed scheduling policy (sched_setscheduler)
 - Migrating tasks between CPUs
 - Moving tasks between groups
- Flags
 - #define ENQUEUE_WAKEUP 1
 - #define ENQUEUE_HEAD 2
 - #define DEQUEUE_SLEEP 1

Scheduling classes

- `yield_task`
 - Called as a result of `sched_yield` system call by the running task
 - Yields CPU to other tasks
- `check_preempt_curr`
 - Should check if the current task should be preempted by a new ready task and call `resched_task` if so
 - Called on a (new) task wakeup

Scheduling classes

- `pick_next_task`
 - Should pick the next task to execute
 - Called from `schedule()` to get the task to switch to
- `put_prev_task`
 - Called when the running task is rescheduled
 - Called from `schedule()` right before `pick_next_task`
- `set_curr_task`
 - Called when task priority, group or scheduling policy changes, if the task is running at the time (`rt_mutex_setprio`, `sched_setscheduler`)
 - Always called in combination with `put_prev_task`

Scheduling classes

- `task_tick`
 - Called on timer with HZ frequency
 - This is the place to do time accounting and to run timeslice-based preemption
- `switched_from/switched_to`
 - Called when the scheduling class changes (due to the change in priority)
- `prio_changed`
 - Called when the priority changes, but not the scheduling class (`rt_mutex_setprio`, `sched_setscheduler`)

Assignment: Dummy scheduler

- FCFS non-preemptive scheduler with one priority level
- Your task:
 - Priority scheduling with support for 5 priority levels
 - Preemption due to a task of a higher priority becoming available
 - Preemption due to running task's timeslice expiry
 - A mechanism to prevent the starvation of processes with lower priority (ageing)
- To have a task scheduled by the dummy scheduler, you need to assign it a priority in the range [15-19]

UML

- User Mode Linux is a Linux virtual machine that runs on a Linux host
- It is included in the mainline kernel
- Download sources from kernel.org and build using

```
make ARCH=um defconfig
make ARCH=um
```

References

- Understanding the Linux Kernel, Daniel Bovet, Marco Cesati; O'Reilly, 3rd eddition
- Professional Linux Kernel Architecture, Wolfgang Mauerer